

inciliardward/Softward interface

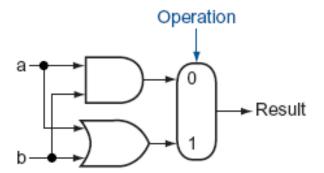
# **Chapter 3**

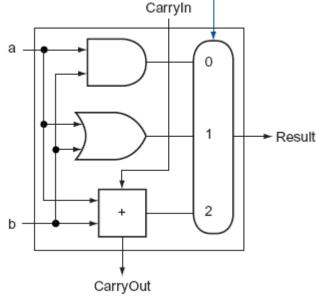
# **Arithmetic for Computers**

# **Basic Arithmetic Logic Unit**

One-bit ALU that performs AND, OR, and

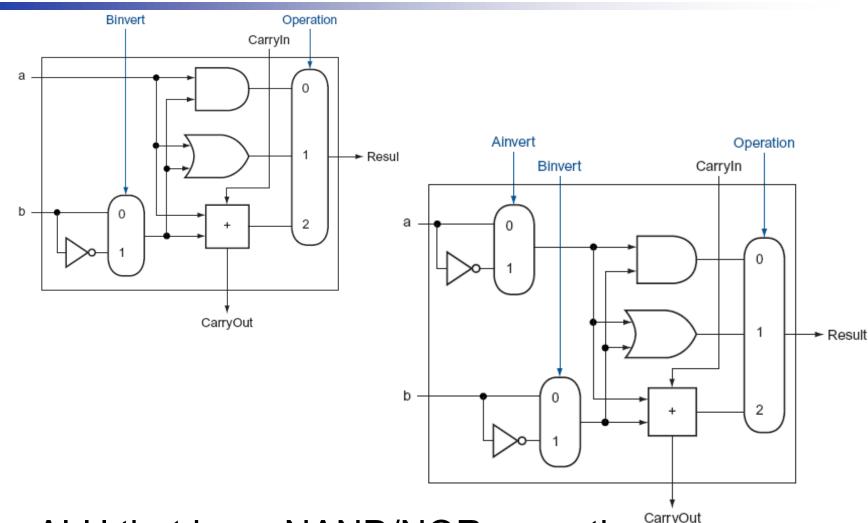
addition







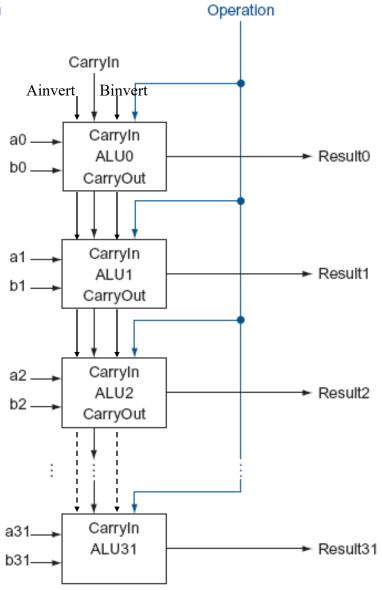
### **Enhanced Arithmetic Logic Unit**



ALU that have NAND/NOR operation Carryout



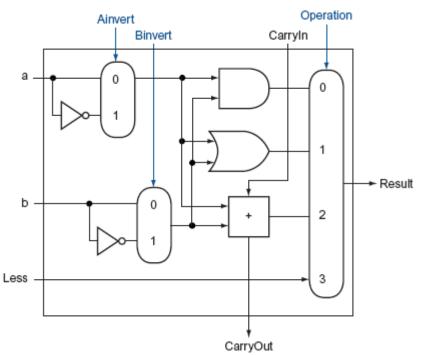
### 32-bit ALU

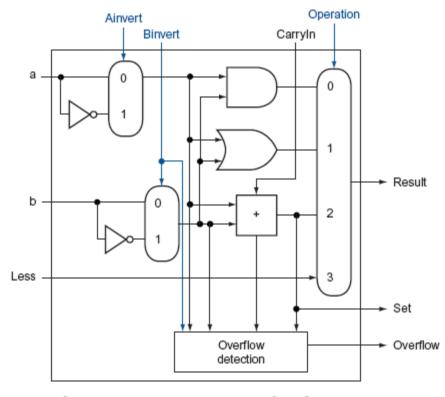




#### **One-bit ALUs with Set Less Than**

- Set less than instruction requires a subtraction and then sets all but the least significant bit to 0, with the lsb set to 1 if a < b</li>
- Less signal line
  - Isb signed bit
  - All but the lsb 0



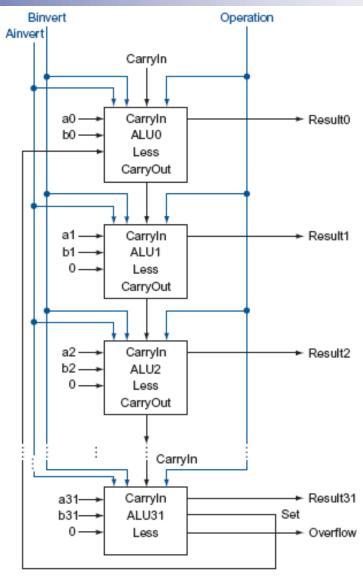




### 32-bit ALU with Set Less Than

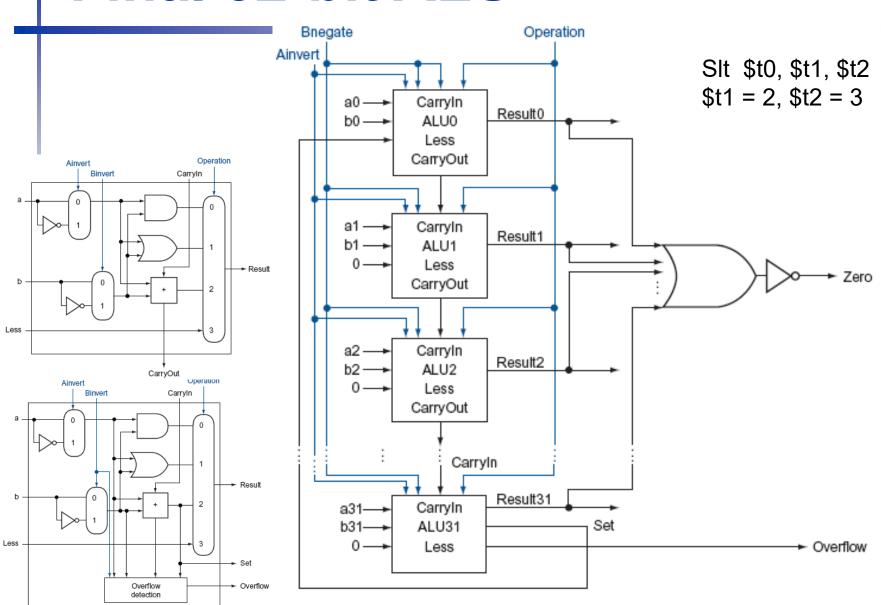
If a < b output 1 else output 0

000...00





### Final 32-bit ALU



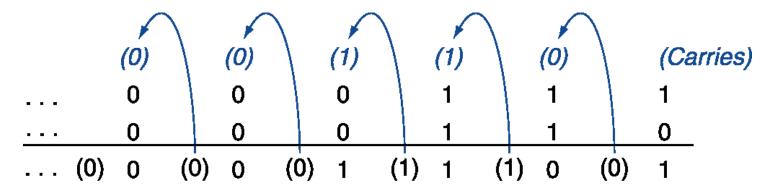
# **Arithmetic for Computers**

- Operations on integers
  - Addition and subtraction
  - Multiplication and division
  - Dealing with overflow
- Floating-point real numbers
  - Representation and operations



# **Integer Addition**

Example: 7 + 6



- Overflow if result out of range
  - Adding +ve and –ve operands, no overflow
  - Adding two +ve operands
    - Overflow if result sign is 1
  - Adding two –ve operands
    - Overflow if result sign is 0



# Integer Subtraction

- Add negation of second operand
- Example: 7 6 = 7 + (-6)

```
+7: 0000 0000 ... 0000 0111
```

- +1: 0000 0000 ... 0000 0001
- Overflow if result out of range
  - Subtracting two +ve or two –ve operands, no overflow
  - Subtracting +ve from –ve operand (ex. (-b) a)
    - Overflow if result sign is 0
  - Subtracting –ve from +ve operand (ex. b (-a))
    - Overflow if result sign is 1



# **Dealing with Overflow**

- Some languages (e.g., C) ignore overflow
  - Use MIPS addu, addui, subu instructions
- Other languages (e.g., Ada, Fortran) require raising an exception
  - Use MIPS add, addi, sub instructions
  - On overflow, invoke exception handler
    - Save PC in exception program counter (EPC) register
    - Jump to predefined handler address
    - mfc0 (move from coprocessor reg) instruction can retrieve EPC value, to return after corrective action



# Overflow Detection for Signed & Unsigned Addition

Signed addition addu \$t0, \$t1, \$t2

xor \$t3, \$t1, \$t2

slt \$t3, \$t3, \$zero

bne \$t3, \$zero, No\_overflow

xor \$t3, \$t0, \$t1

slt \$t3, \$t3, \$zero

bne \$t3, \$zero, Overflow

Unsigned addition

addu \$t0, \$t1, \$t2

nor \$t3, \$t1, \$zero

sltu \$t3, \$t3, \$t2

bne \$t3, \$zero, Overflow

01101

10010



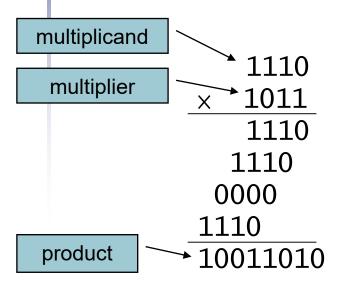
### **Arithmetic for Multimedia**

- Graphics and media processing operates on vectors of 8-bit and 16-bit data
  - Use 64-bit adder, with partitioned carry chain
    - Operate on 8x8-bit, 4x16-bit, or 2x32-bit vectors
  - SIMD (single-instruction, multiple-data)
- Saturating operations
  - On overflow, result is largest representable value
    - c.f. 2s-complement modulo arithmetic
  - E.g., clipping in audio, saturation in video

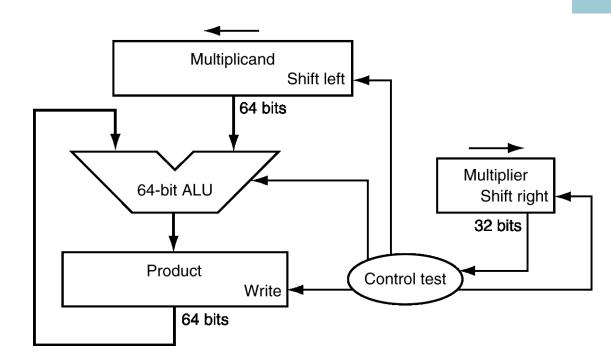


# Multiplication

Start with long-multiplication approach

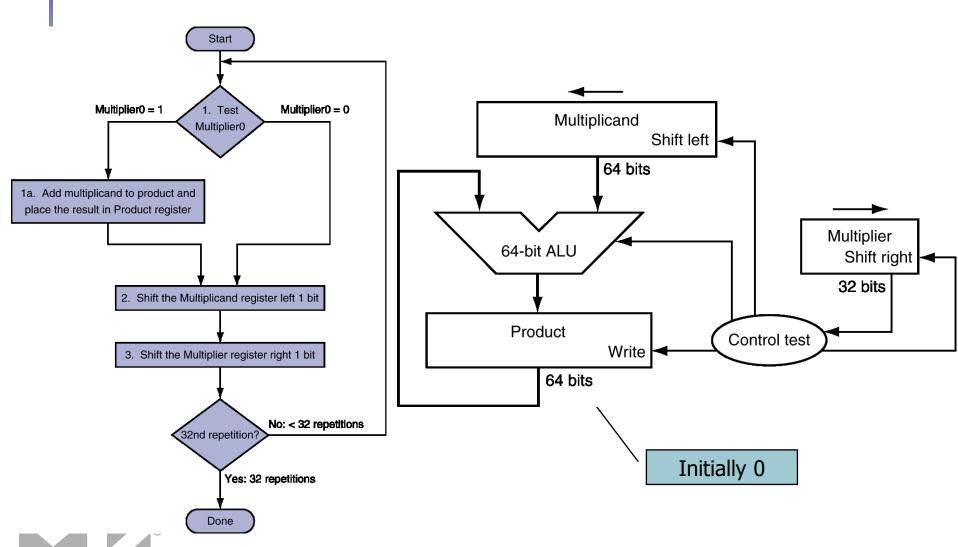


Length of product is the sum of operand lengths





# **Multiplication Hardware**







1011

add

shift

add

shift

shift

1000 **1011** Perform steps in parallel: add/shift 0010 0111

01000 **101** 1000

11000 **101** 

011000 10 0000

add

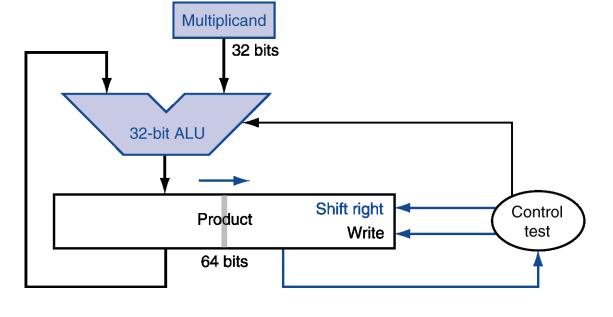
0110<mark>00 **10**</mark> 0011<mark>000 **1**</mark>

1000

add

1011<mark>000 1</mark>

shift 01011000



00110 000110 0010

001110

0010

0010

00110

0010

001110

0010

0111

0010

0010

00010

One cycle per partial-product addition

That's ok, if frequency of multiplications is low

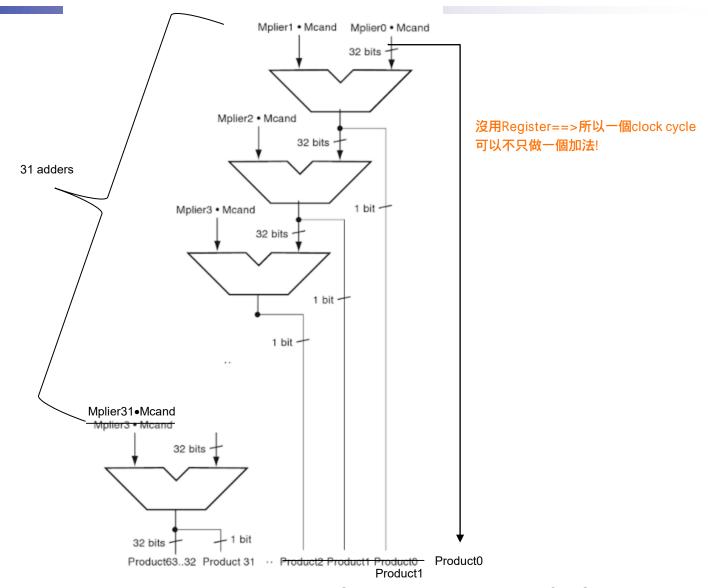


# **Multiplication Example**

Iteration	Step	Product Register		Multiplicand
		(Product	: Multiplier)	
0	Initial value	0000	0011	0010
1	1: 1→Prod+=Mcand	0010	0011	0010
	2: shift right Preg	0001	0001	0010
2	1: 1→Prod+=Mcand	0011	0001	0010
	2: shift right Preg	0001	1000	0010
3	1: 0→no operation	0001	1000	0010
	2: shift right Preg	0000	1100	0010
4	1: 0→no operation	0000	1100	0010
	2: shift right Preg	0000	0110	0010



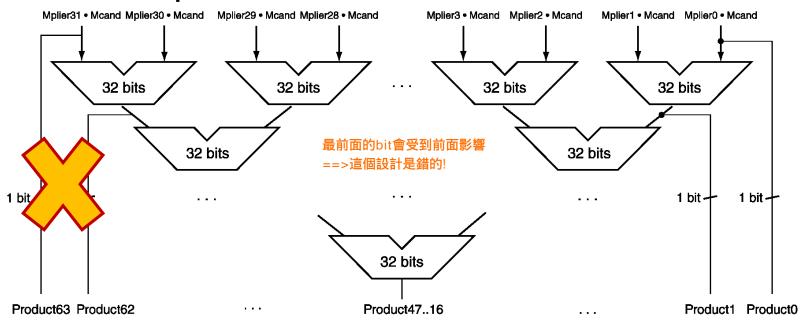
# **Faster Multiplier**





# **Faster Multiplier**

- Uses multiple adders
  - Cost/performance tradeoff



- Can be pipelined
  - Several multiplication performed in parallel



# Fast Carry Using the First Level of Abstraction

◆ ci+1: carry output of level i, carry input of level i+1

$$ci+1 = (bi \cdot ci) + (ai \cdot ci) + (ai \cdot bi)$$
$$= (ai \cdot bi) + (ai + bi) \cdot ci$$

• For example

$$c2 = (a1 \cdot b1) + (a1 + b1) \cdot ((a0 \cdot b0) + (a0 + b0) \cdot c0)$$

- We can define generate gi and propagate pi  $gi = ai \cdot bi$ such that  $ci+1 = gi + pi \cdot ci$
- if ai = bi = 1  $ci+1 = gi + pi \cdot ci = 1 + pi \cdot ci = 1$
- $if \ ai = 1, \ bi = 0 \ or \ ai = 0, \ bi = 1$   $ci+1 = gi + pi \cdot ci = 0 + 1 \cdot ci = ci$
- $c1 = g0 + (p0 \cdot c0)$   $c2 = g1 + (p1 \cdot c1) = g1 + (p1 \cdot g0) + (p1 \cdot p0 \cdot c0)$   $c3 = g2 + (p2 \cdot c2) = g2 + (p2 \cdot g1) + (p2 \cdot p1 \cdot g0) + (p2 \cdot p1 \cdot p0 \cdot c0)$   $c4 = g3 + (p3 \cdot g2) + (p3 \cdot p2 \cdot g1) + (p3 \cdot p2 \cdot p1 \cdot g0)$  $c4 + (p3 \cdot p2 \cdot p1 \cdot p0 \cdot c0)$

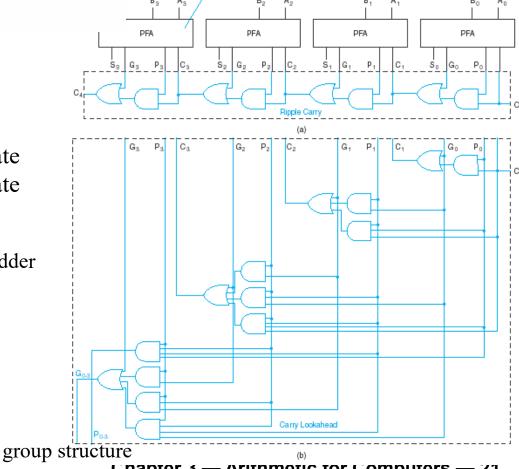


# 4-bit Carry Look-Ahead Adder

```
c1 = g0 + (p0 \cdot c0)
c2 = g1 + (p1 \cdot c1) = g1 + (p1 \cdot g0) + (p1 \cdot p0 \cdot c0)
c3 = g2 + (p2 \cdot c2) = g2 + (p2 \cdot g1) + (p2 \cdot p1 \cdot g0) + (p2 \cdot p1 \cdot p0 \cdot c0)
c4 = g3 + (p3 \cdot g2) + (p3 \cdot p2 \cdot g1) + (p3 \cdot p2 \cdot p1 \cdot g0)
+ (p3 \cdot p2 \cdot p1 \cdot p0 \cdot c0)
```

 $C_4$ : 4 AND gates and 1 OR gate  $C_n$ : n AND gates and 1 OR gate

4-bit CLA adder





# Fast Carry Using the Second Level of Abstraction

The concept can be extended to another level by considering *group generate* (g0-3) and *group propagate* (p0-3) functions:

$$g0-3 = g3 + p3g2 + p3p2g1 + p3p2p1g0$$

$$p0-3 = p3p2p1p0$$

$$C_4 = g_3 + p_3 c_3 = g_3 + p_3g_2 + p_3p_2g_1 + p_3p_2p_1g_0$$

$$+ p_3p_3p_3 + p_3p_3p_3 + p_3p_3p_3p_3 + p_3p_3p_3p_3p_3p_3$$

Using these two equations:

$$c4 = g0-3 + p0-3c0$$

$$c8 = g4-7 + p4-7c4$$
 $= g4-7 + p4-7(g0-3 + p0-3c0)_{\text{Level 2}}$ 
 $= g4-7 + p4-7g0-3 + p4-7p0-3c0$ 

 Thus, it is possible to have four 4-bit adders that use one of the same carry lookahead circuit to speed up 16-bit addition



0001 1110 1010 1011

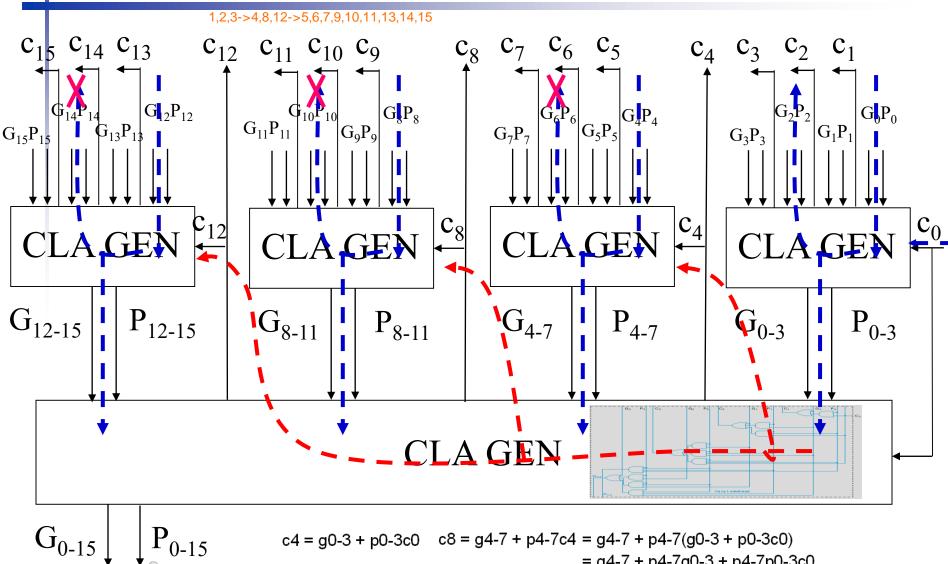
Carry lookahead circuitry

1001 0010 1111 0011

|0001||1110||1010||101

Carry lookahead circuitry

### 16-bit Two-Level Carry Look-Ahead **Adder**



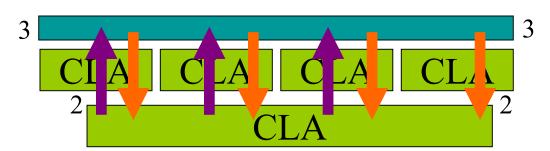
= q4-7 + p4-7q0-3 + p4-7p0-3c0

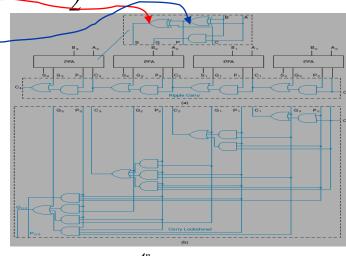
Chapter 3 — Arithmetic for Computers — 23

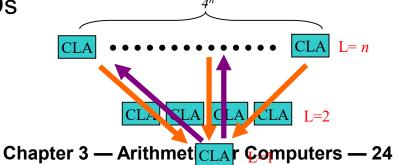
### **Carry Lookahead Example**

#### Specifications 1:

- 16-bit CLA
- Delays:
  - NOT = 1
  - XOR = Isolated AND = 3
  - AND-OR = 2
- Longest Delays:
  - Ripple carry adder\* = 3 + 15 × 2 + 3 = 36
  - CLA =  $3 + 3 \times 2 + 3 = 12$
- Specification 2:
  - Exclusive OR = 2 gate delays (GDs)
  - 2-level 16-bit CLA delay = 10 GDs
  - 3-level 64-bit CLA delay = 14 GDs
  - n-level  $4^n$ -bit CLA delay = 4n + 2









# Simplified Multiplication

- Consider 01110 =  $1 \times 2^3 + 1 \times 2^2 + 1 \times 2^1$  (three additions)
- One important observation another faster calculation
  - $01110 = 1 \times 2^4 1 \times 2^1$  (one addition and one subtraction)
- Multiplication has similar property
  - The process on the left is traditional operation
  - The process on the right applies the above concept

```
 0010 \times 0110 = 0 \times (0010 \times 2^0) + 1 \times (0010 \times 2^1) + 1 \times (0010 \times 2^2) + 0 \times (0010 \times 2^3)
```

$$0010 \times 0110 = 0 \times (0010 \times 2^{0}) - 1 \times (0010 \times 2^{1}) + 0 \times (0010 \times 2^{2}) + 1 \times (0010 \times 2^{3})$$

```
\begin{array}{c} 0010_{\text{two}} \\ \times & 0110_{\text{two}} \\ + & 0000 \text{ shift (0 in multiplier)} \\ + & 0010 \text{ add (1 in multiplier)} \\ + & 0010 \text{ add (1 in multiplier)} \\ + & 0000 \text{ shift (0 in multiplier)} \\ \hline \hline 00001100_{\text{two}} \end{array}
```

```
\begin{array}{c} 0010_{\text{two}} \\ \times & 0110_{\text{two}} \\ + & 0000 \quad \text{shift (0 in multiplier)} \\ - & 0010 \quad \text{sub (first 1 in multiplier)} \\ + & 0000 \quad \quad \text{shift (middle of string of 1s)} \\ + & 0010 \quad \quad \text{add (prior step had last 1)} \\ 00001100_{\text{two}} \end{array}
```



### **Booth's Algorithm**

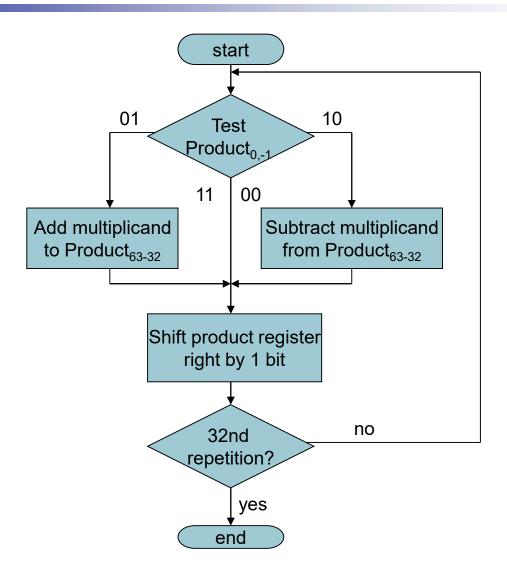
Current bit	Bit to the right	Explanation	example
1	0	Beginning of a run of 1s	00001111000
1	1	Middle of a run of 1s	00001111000
0	1	End of a run of 1s	000 <mark>01</mark> 111000
0	0	Middle of a run of 0s	00001111000

#### Booth's algorithm

- Based on the current and previous bits, do one of the following
  - 00: middle of a string of 0s, so no arithmetic operation.
  - 01: end of a string of 1s, so add the multiplicand to the left half of the product
  - 10: beginning of a string of 1s, so subtract the multiplicand from the left half of the product.
  - 11: middle of a string of 1s, so no arithmetic operation.
- As in the previous algorithm, shift the product register right 1 bit



# **Booth's Algorithm**





# **Examples for Booth's Algorithm**

Itera- Multi-		Original algorithm		Booth's algorithm	
tion plica	plicand	Step	Product	Step	Product
0	0010	Initial values	0000 0110	Initial values	0000 0110 0
1	0010	1: 0 ⇒ no operation	0000 0110	1a: 00 ⇒ no operation	0000 0110 0
1	0010	2: Shift right Product	0000 0011	2: Shift right Product	0000 0011 0
2	0010	1a: 1 ⇒ Prod = Prod + Mcand	0010 0011	1c: 10 ⇒ Prod = Prod - Mcand	1110 0011 0
	0010	2: Shift right Product	0001 0001	2: Shift right Product	1111 0001 1
3	0010	1a: 1 ⇒ Prod = Prod + Mcand	0011 0001	1d: 11 ⇒ no operation	1111 0001 1
	0010	2: Shift right Product	0001 1000	2: Shift right Product	1111 1000 1
4	0010	1: 0 ⇒ no operation	0001 1000	1b: 01 ⇒ Prod = Prod + Mcand	0001 1000 1
	0010	2: Shift right Product	0000 1100	2: Shift right Product	0000 1100 0

Iteration	Step	Multiplicand	Product
0	Initial values	0010	0000 1101 0
1	1c: 10 ⇒ Prod = Prod - Mcand	0010	1110 1101 0
	2: Shift right Product	0010	1111 0110 1
2	2 1b: 01 ⇒ Prod = Prod + Mcand		0001 0110 1
	2: Shift right Product	0010	0000 1011 0
3	3 1c: 10 ⇒ Prod = Prod – Mcand		1110 1011 0
	2: Shift right Product	0010	1111 0101 1
4	4 1d: 11 ⇒ no operation		1111 0101 1
	2: Shift right Product	0010	1111 1010 1

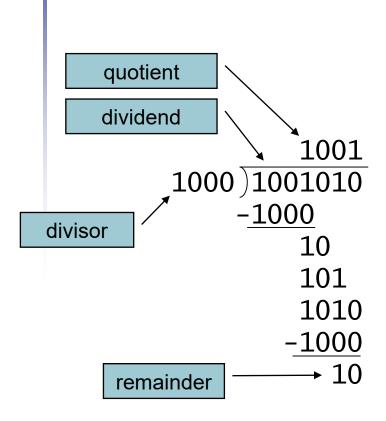


# **MIPS Multiplication**

- Two 32-bit registers for product
  - HI: most-significant 32 bits
  - LO: least-significant 32 bits
- Instructions
  - mult rs, rt / multu rs, rt
    - 64-bit product in HI/LO
  - mfhi rd / mflo rd
    - Move from HI/LO to rd
    - Can test HI value to see if product overflows 32 bits
  - mul rd, rs, rt
    - Least-significant 32 bits of product —> rd



### **Division**

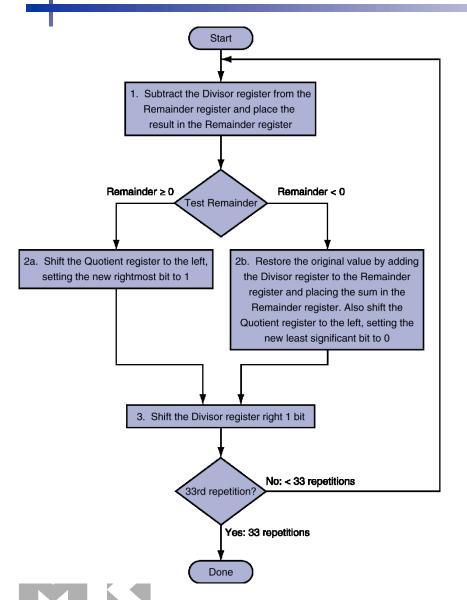


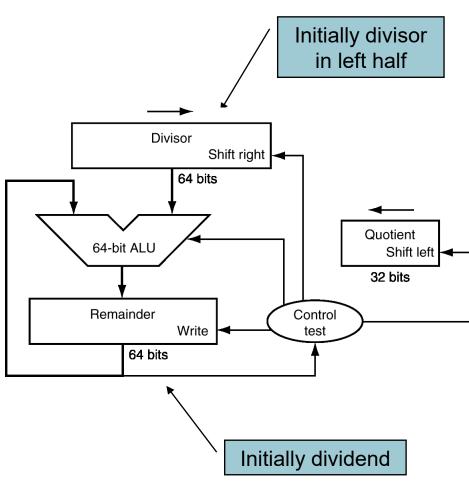
*n*-bit operands yield *n*-bit quotient and remainder

- Check for 0 divisor
- Long division approach
  - If divisor ≤ dividend bits
    - 1 bit in quotient, subtract
  - Otherwise
    - 0 bit in quotient, bring down next dividend bit
- Restoring division
  - Do the subtract, and if remainder goes < 0, add divisor back</li>
- Signed division
  - Divide using absolute values
  - Adjust sign of quotient and remainder as required



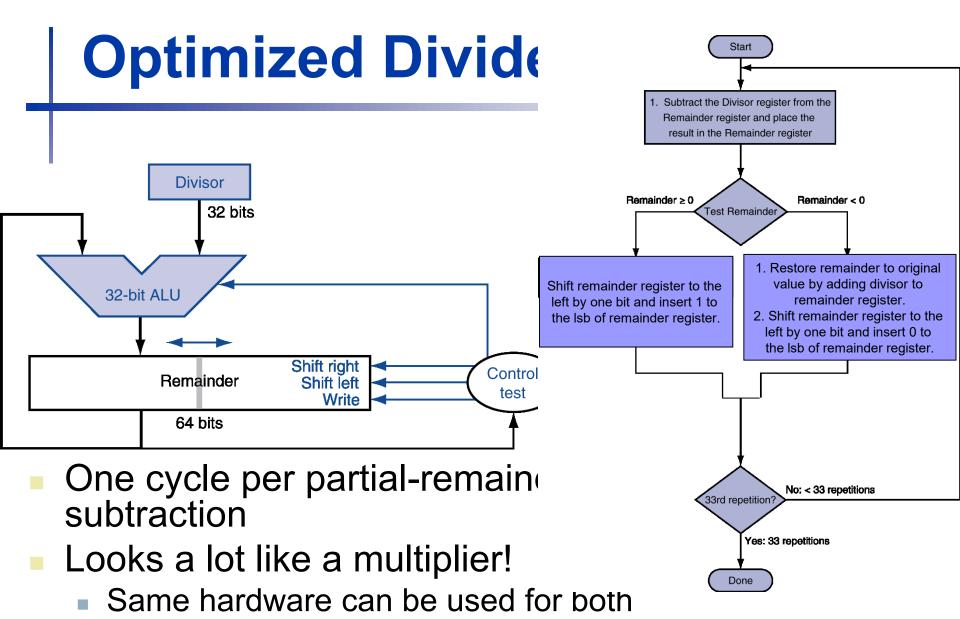
### **Division Hardware**





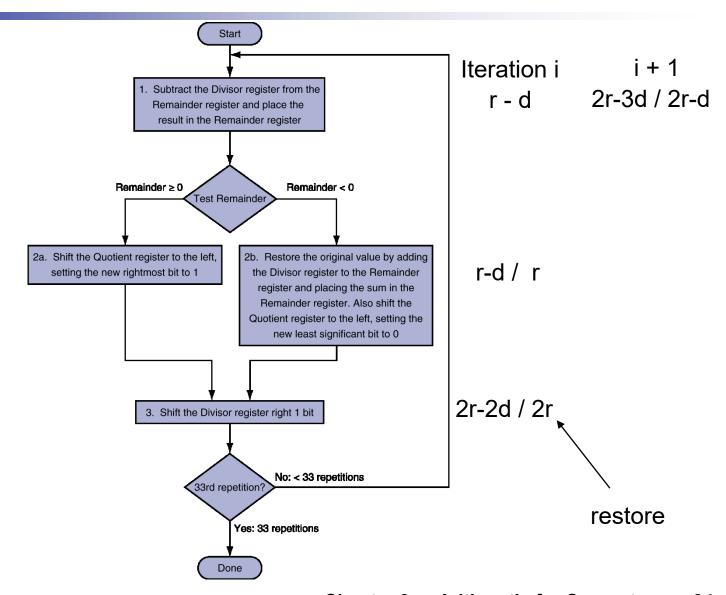
# **Division Example**

Iteration	Step	Quotient	Divisor	Remainder
0	Initial values	0000	0010 0000	0000 0111
	1: Rem = Rem - Div	0000	0010 0000	<b>1</b> 110 0111
1	2b: Rem < $0 \Rightarrow$ +Div, sll Q, QQ = $0$	0000	0010 0000	0000 0111
	3: Shift Div right	0000	0001 0000	0000 0111
2	1: Rem = Rem - Div	0000	0001 0000	<mark>1</mark> 111 0111
	2b: Rem < $0 \Rightarrow$ +Div, sll Q, QQ = $0$	0000	0001 0000	0000 0111
	3: Shift Div right	0000	0000 1000	0000 0111
3	1: Rem = Rem - Div	0000	0000 1000	<mark>1</mark> 111 1111
	2b: Rem < $0 \Rightarrow$ +Div, sll Q, QQ = $0$	0000	0000 1000	0000 0111
	3: Shift Div right	0000	0000 0100	0000 0111
4	1: Rem = Rem - Div	0000	0000 0100	0000 0011
	2a: Rem ≥ 0 ⇒ sll Q, QQ = 1	0001	0000 0100	0000 0011
	3: Shift Div right	0001	0000 0010	0000 0011
5	1: Rem = Rem – Div	0001	0000 0010	0000 0001
	2a: Rem ≥ 0 ⇒ sll Q, QQ = 1	0011	0000 0010	0000 0001
	3: Shift Div right	0011	0000 0001	0000 0001





# **Restoring Algorithm**





**Division Algorithm** 

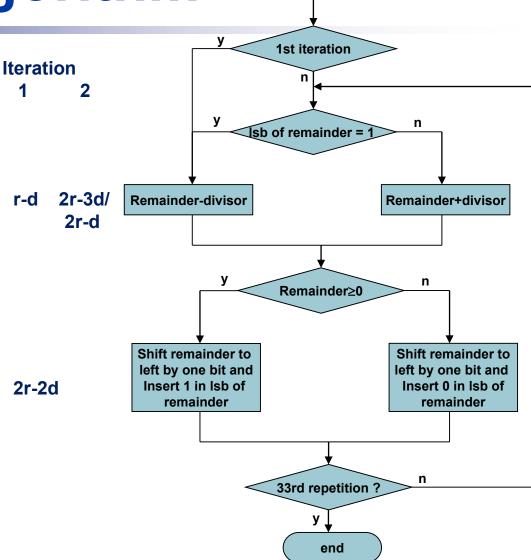
- Restoring algorithm
  - $r d \rightarrow quotient = 1$
  - sll: 2r → quotient = 0
     next

iteration: 2r – d (two subtractions and one addition)

- Non-restoring algorithm
  - $r d \rightarrow quotient = 1$
  - sll: 2(r d) → quotient = 0

next iteration: 2(r-d) + d = 2r - d (one subtraction and addition)

 non-restoring flow is in the right





### **Faster Division**

- Can't use parallel hardware as in multiplier
  - Subtraction is conditional on sign of remainder
- Faster dividers (e.g. SRT devision)
   generate multiple quotient bits per step
  - Still require multiple steps



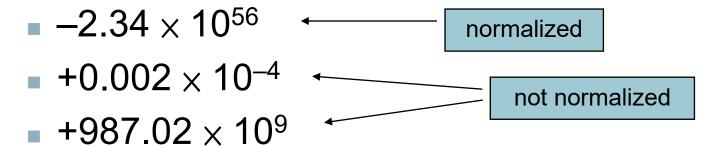
#### **MIPS Division**

- Use HI/LO registers for result
  - HI: 32-bit remainder
  - LO: 32-bit quotient
- Instructions
  - div rs, rt / divu rs, rt
  - No overflow or divide-by-0 checking
    - Software must perform checks if required
  - Use mfhi, mflo to access result



# **Floating Point**

- Representation for non-integral numbers
  - Including very small and very large numbers
- Like scientific notation



- In binary
  - $\pm 1.xxxxxxx_2 \times 2^{yyyy}$
- Types float and double in C



# Floating Point Standard

- Defined by IEEE Std 754-1985
- Developed in response to divergence of representations
  - Portability issues for scientific code
- Now almost universally adopted
- Two representations
  - Single precision (32-bit)
  - Double precision (64-bit)



# **IEEE Floating-Point Format**

single: 8 bits single: 23 bits double: 11 bits double: 52 bits

S Exponent Fraction

$$x = (-1)^{S} \times (1 + Fraction) \times 2^{(Exponent-Bias)}$$

- S: sign bit  $(0 \Rightarrow \text{non-negative}, 1 \Rightarrow \text{negative})$
- Normalize significand: 1.0 ≤ |significand| < 2.0</p>
  - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
  - Significand is Fraction with the "1." restored
- Exponent: excess representation: actual exponent + Bias
  - Ensures exponent is unsigned
  - Single: Bias = 127; Double: Bias = 1023



# Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
  - Exponent: 00000001⇒ actual exponent = 1 - 127 = -126
  - Fraction:  $000...00 \Rightarrow \text{significand} = 1.0$
  - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
  - exponent: 11111110⇒ actual exponent = 254 127 = +127
  - Fraction: 111...11 ⇒ significand ≈ 2.0
  - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

8 bits	127	254	
	0	127	
	-1	126	
	 -126	1	
		^	
	-127	0	
	-128	255	



## **Double-Precision Range**

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
  - Exponent: 0000000001⇒ actual exponent = 1 - 1023 = -1022
  - Fraction:  $000...00 \Rightarrow \text{significand} = 1.0$
  - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
  - Exponent: 11111111110⇒ actual exponent = 2046 1023 = +1023
  - Fraction: 111...11 ⇒ significand ≈ 2.0
  - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$



## Floating-Point Precision

- Relative precision
  - all fraction bits are significant
  - Single: approx 2<sup>-23</sup>
    - Equivalent to 23 x log<sub>10</sub>2 ≈ 23 x 0.3 ≈ 6 decimal digits of precision
  - Double: approx 2<sup>-52</sup>
    - Equivalent to 52 x log<sub>10</sub>2 ≈ 52 x 0.3 ≈ 16 decimal digits of precision



# Floating-Point Example

- Represent –0.75
  - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
  - S = 1
  - Fraction =  $1000...00_2$
  - Exponent = -1 + Bias
    - Single:  $-1 + 127 = 126 = 011111110_2$
    - Double:  $-1 + 1023 = 1022 = 0111111111110_2$
- Single: 1011111101000...00
- Double: 10111111111101000...00



## Floating-Point Example

 What number is represented by the singleprecision float

11000000101000...00

- S = 1
- Fraction =  $01000...00_2$
- Exponent =  $10000001_2 = 129$

$$x = (-1)^{1} \times (1 + 01_{2}) \times 2^{(129 - 127)}$$

$$= (-1) \times 1.25 \times 2^{2}$$

$$= -5.0$$



#### **Denormal Numbers**

Exponent = 000...0 with fraction ≠ 0 ⇒ hidden bit is 0

$$x = (-1)^{S} \times (0 + Fraction) \times 2^{1-Bias}$$

- Smaller than normal numbers
  - allow for gradual underflow, with diminishing precision
- Denormal with fraction = 000...0

$$x = (-1)^{S} \times (0+0) \times 2^{-Bias} = \pm 0.0$$



Two representations of 0.0!

### **Infinities and NaNs**

- Exponent = 111...1, Fraction = 000...0
  - ±Infinity
  - Can be used in subsequent calculations, avoiding need for overflow check
- Exponent = 111...1, Fraction ≠ 000...0
  - Not-a-Number (NaN)
  - Indicates illegal or undefined result
    - e.g., 0.0 / 0.0
  - Can be used in subsequent calculations



# **IEEE 754 Encoding of FPN**

Single F	Precision	Double Precision		Object represented	
Exponent	Fraction	Exponent	Fraction		
0	0	0	0	0	
0	Nonzero	0	Nonzero	±denormalized number	
1-254	Anything	1-2046	Anything	±floating-point number	
255	0	2047	0	± ∞	
255	Nonzero	2047	Nonzero	NaN	

- Smallest positive single precision normalized number
- Smallest positive single precision denormalized no. (Hint: Fraction is 23-bit)
   Exponent:-126 fraction:22個0搭配最後一個1
- $\infty$  must obey mathematical conventions: F +  $\infty$  =  $\infty$ ; F/  $\infty$  = 0



## Floating-Point Addition

- Consider a 4-digit decimal example
  - $-9.999 \times 10^{1} + 1.610 \times 10^{-1}$
- 1. Align decimal points
  - Shift number with smaller exponent
  - $\bullet$  9.999 × 10<sup>1</sup> + 0.016 × 10<sup>1</sup>
- 2. Add significands
  - $\bullet$  9.999 × 10<sup>1</sup> + 0.016 × 10<sup>1</sup> = 10.015 × 10<sup>1</sup>
- 3. Normalize result & check for over/underflow
  - $\blacksquare$  1.0015  $\times$  10<sup>2</sup>
- 4. Round and renormalize if necessary
  - $1.002 \times 10^2$



# Floating-Point Addition

- Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2} (0.5 + -0.4375)$
- 1. Align binary points
  - Shift number with smaller exponent
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
- 2. Add significands
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
- 3. Normalize result & check for over/underflow
  - $1.000_2 \times 2^{-4}$ , with no over/underflow
- 4. Round and renormalize if necessary
  - $-1.000_2 \times 2^{-4}$  (no change) = 0.0625

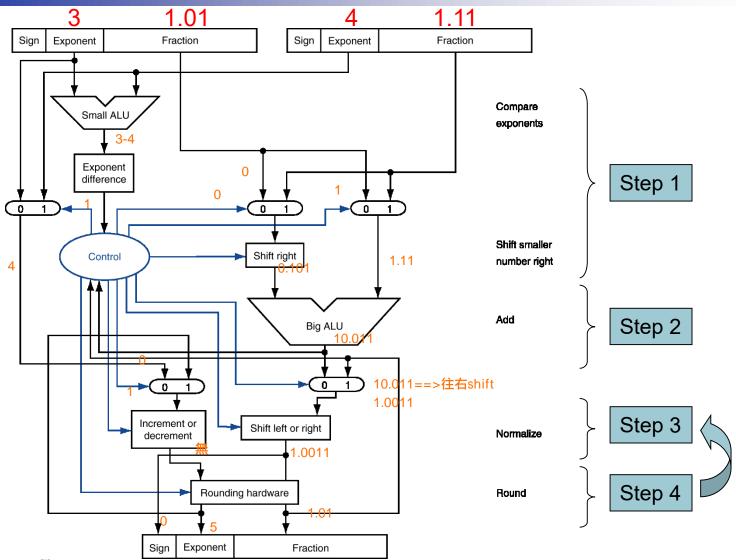


#### **FP Adder Hardware**

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
  - Much longer than integer operations
  - Slower clock would penalize all instructions
- FP adder usually takes several cycles
  - Can be pipelined



### **FP Adder Hardware**





# Floating-Point Multiplication

- Consider a 4-digit decimal example
  - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- 1. Add exponents
  - For biased exponents, subtract bias from sum
  - New exponent = 10 + -5 = 5
- 2. Multiply significands
  - $1.110 \times 9.200 = 10.212 \implies 10.212 \times 10^{5}$
- 3. Normalize result & check for over/underflow
  - $\bullet$  1.0212 × 10<sup>6</sup>
- 4. Round and renormalize if necessary
  - 1.021 × 10<sup>6</sup>
- 5. Determine sign of result from signs of operands
  - $+1.021 \times 10^6$



# Floating-Point Multiplication

- Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2} (0.5 \times -0.4375)$
- 1. Add exponents
  - Unbiased: -1 + -2 = -3
  - Biased: (-1 + 127) + (-2 + 127) = -3 + 254 127 = -3 + 127
- 2. Multiply significands
  - $1.000_2 \times 1.110_2 = 1.110_2 \implies 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
  - $1.110_2 \times 2^{-3}$  (no change) with no over/underflow
- 4. Round and renormalize if necessary
  - $1.110_2 \times 2^{-3}$  (no change)
- 5. Determine sign: +ve  $\times$  –ve  $\Rightarrow$  –ve
  - $-1.110_2 \times 2^{-3} = -0.21875$



### **FP Arithmetic Hardware**

- FP multiplier is of similar complexity to FP adder
  - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
  - Addition, subtraction, multiplication, division, reciprocal, square-root
  - FP ↔ integer conversion
- Operations usually takes several cycles
  - Can be pipelined



### **FP Instructions in MIPS**

- FP hardware is coprocessor 1
  - Adjunct processor that extends the ISA
- Separate FP registers
  - 32 single-precision: \$f0, \$f1, ... \$f31
  - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
    - Release 2 of MIPs ISA supports 32 x 64-bit FP reg's
- FP instructions operate only on FP registers
  - Programs generally don't do integer ops on FP data, or vice versa
  - More registers with minimal code-size impact
- FP load and store instructions
  - lwc1, ldc1, swc1, sdc1
    - e.g., ldc1 \$f8, 32(\$sp)



#### **FP Instructions in MIPS**

- Single-precision arithmetic
  - add.s, sub.s, mul.s, div.s
    - e.g., add.s \$f0, \$f1, \$f6
- Double-precision arithmetic
  - add.d, sub.d, mul.d, div.d
    - e.g., mul.d \$f4, \$f4, \$f6
- Single- and double-precision comparison
  - c.xx.s, c.xx.d (xx is eq, lt, le, ...)
  - Sets or clears FP condition-code bit
    - e.g. c.lt.s \$f3, \$f4
- Branch on FP condition code true or false
  - bc1t, bc1f
    - e.g., bc1t 25



# FP Example: °F to °C

C code:

```
float f2c (float fahr) {
  return ((5.0/9.0)*(fahr - 32.0));
}
```

- fahr in \$f12, result in \$f0, literals in global memory space
- Compiled MIPS code:

```
f2c: lwc1  $f16, const5($gp)
  lwc2  $f18, const9($gp)
  div.s  $f16, $f16, $f18
  lwc1  $f18, const32($gp)
  sub.s  $f18, $f12, $f18
  mul.s  $f0, $f16, $f18
  jr  $ra
```



### FP Example: Array Multiplication

- $X = X + Y \times Z$ 
  - All 32 x 32 matrices, 64-bit double-precision elements
- C code:

Addresses of x, y, z in \$a0, \$a1, \$a2, and i, j, k in \$s0, \$s1, \$s2



### FP Example: Array Multiplication

#### MIPS code:

```
li $t1, 32
                   # t1 = 32 (row size/loop end)
   li $s0, 0
                   # i = 0; initialize 1st for loop
L1: li $s1, 0
                   # j = 0; restart 2nd for loop
L2: 1i $s2, 0 # k = 0; restart 3rd for loop
   addu t2, t2, t2, t2 = i * size(row) + j
   sll $t2, $t2, 3 # $t2 = byte offset of [i][j]
   addu t2, a0, t2 \# t2 = byte address of <math>x[i][j]
   1.d f4, 0(f2) # f4 = 8 bytes of x[i][j]
L3: s11 $t0, $s2, 5 # $t0 = k * 32 (size of row of z)
   addu t0, t0, s1 # t0 = k * size(row) + j
   sll $t0, $t0, 3 # $t0 = byte offset of [k][j]
   addu t0, a2, t0 # t0 = byte address of <math>z[k][j]
   1.d f16, 0(t0) # f16 = 8 bytes of z[k][j]
```



### FP Example: Array Multiplication

•••

```
\$11 \$t0, \$s0, 5  # \$t0 = i*32 (size of row of y)
addu $t0, $t0, $s2  # $t0 = i*size(row) + k
sll $t0, $t0, 3 # $t0 = byte offset of [i][k]
addu $t0, $a1, $t0  # $t0 = byte address of y[i][k]
1.d f18, 0(t0) # f18 = 8 bytes of y[i][k]
mul.d f16, f18, f16 # f16 = y[i][k] * z[k][j]
add.d f4, f4, f4 # f4=x[i][j] + y[i][k]*z[k][j]
addiu $s2, $s2, 1 # $k k + 1
bne $s2, $t1, L3 # if (k != 32) go to L3
s.d f4, 0(t2) # x[i][j] = f4
addiu $s1, $s1, 1 # $j = j + 1
bne $s1, $t1, L2 # if (j != 32) go to L2
addiu $s0, $s0, 1
                    # $i = i + 1
bne $s0, $t1, L1 # if (i != 32) go to L1
```



### **Accuracy of Floating-Point Operations**

- Consider the addition:  $2.56 \times 10^{0} + 2.34 \times 10^{2}$   $0.02\underline{56} \times 10^{2} + 2.34\underline{00} \times 10^{2} = 2.3656 \times 10^{2} = 2.37 \times 10^{2}$ Guard and round O.44 ulp (unit in the last place)
- If there are no guard and round extra bits, the result will be  $0.02 \times 10^2 + 2.34 \times 10^2 = 2.36 \times 10^2$
- 2.345000000000 (2.34) vs. 2.345000000001 (2.35) by sticky bit (how to get?)



### Interpretation of Data

#### **The BIG Picture**

- Bits have no inherent meaning
  - Interpretation depends on the instructions applied
- Computer representations of numbers
  - Finite range and precision
  - Need to account for this in programs



# **Associativity**

- Parallel programs may interleave operations in unexpected orders
  - Assumptions of associativity may fail

		(x+y)+z	x+(y+z)
X	-1.50E+38		-1.50E+38
у	1.50E+38	0.00E+00	
Z	1.0	1.0	1.50E+38
		1.00E+00	0.00E+00

 Need to validate parallel programs under varying degrees of parallelism



### x86 FP Architecture

- Originally based on 8087 FP coprocessor
  - 8 x 80-bit extended-precision registers
  - Used as a push-down stack
  - Registers indexed from TOS: ST(0), ST(1), ...
- FP values are 32-bit or 64 in memory
  - Converted on load/store of memory operand
  - Integer operands can also be converted on load/store
- Very difficult to generate and optimize code
  - Result: poor FP performance



### **x86 FP Instructions**

Data transfer	Arithmetic	Compare	Transcendental
FILD mem/ST(i) FISTP mem/ST(i) FLDPI FLD1 FLDZ	FIADDP mem/ST(i) FISUBRP mem/ST(i) FIMULP mem/ST(i) FIDIVRP mem/ST(i) FSQRT FABS FRNDINT	FICOMP FIUCOMP FSTSW AX/mem	FPATAN F2XMI FCOS FPTAN FPREM FPSIN FYL2X

#### Optional variations

- I: integer operand
- P: pop operand from stack
- R: reverse operand order
- But not all combinations allowed



### **Streaming SIMD Extension 2 (SSE2)**

- Adds 4 × 128-bit registers
  - Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
  - 2 x 64-bit double precision
  - 4 x 32-bit double precision
  - Instructions operate on them simultaneously
    - Single-Instruction Multiple-Data



### **Right Shift and Division**

- Left shift by i places multiplies an integer by 2<sup>i</sup>
- Right shift divides by 2<sup>i</sup>?
  - Only for unsigned integers
- For signed integers
  - Arithmetic right shift: replicate the sign bit
  - e.g., -5 / 4
    - 11111011<sub>2</sub> >> 2 = 11111110<sub>2</sub> = -2
    - Rounds toward -∞
  - c.f.  $11111011_2 >>> 2 = 001111110_2 = +62$



### Who Cares About FP Accuracy?

- Important for scientific code
  - But for everyday consumer use?
    - "My bank balance is out by 0.0002¢!" ⊗
- The Intel Pentium FDIV bug
  - The market expects accuracy
  - See Colwell, The Pentium Chronicles



# **Concluding Remarks**

- ISAs support arithmetic
  - Signed and unsigned integers
  - Floating-point approximation to reals
- Bounded range and precision
  - Operations can overflow and underflow
- MIPS ISA
  - Core instructions (MIPS cores and MIPS arithmetic cores): 54 most frequently used
    - 100% of SPECINT, 97% of SPECFP
  - Other instructions: less frequent

